



University
of Glasgow | School of
Computing Science

Who's There? Occupancy Prediction with Smart Sensor Boxes

Jamie Sweeney

School of Computing Science
Sir Alwyn Williams Building
University of Glasgow
G12 8QQ

Level 4 Project — September 26, 2017

Abstract

We show how to produce a level 4 project report using latex and pdflatex using the style file l4proj.cls

Education Use Consent

I hereby give my permission for this project to be shown to other University of Glasgow students and to be distributed in an electronic format. **Please note that you are under no obligation to sign this declaration, but doing so would help future students.**

Name: _____ Signature: _____

Contents

1	Introduction	1
1.1	Context	1
1.2	Motivations	1
1.2.1	Primary Motivation	1
1.2.2	Secondary Motivations	2
1.3	Goals	2
2	Background Analysis	3
2.1	Human Detection	3
2.1.1	Infrared Light	3
2.1.2	Digital Images	4
2.1.3	Audio	4
2.1.4	Other	5
2.2	Device Detection	5
2.3	Exitsting products	6
3	Requirements	7
3.1	Functional Requirements	7
3.1.1	Sensor data collection (A)	7
3.1.2	Sensor data processing (B)	7
3.1.3	Reporting (C)	8
3.1.4	Data storage (D)	8
3.1.5	Estimating occupancy (E)	9

3.1.6	Data serving (F)	9
3.2	Non-functional Requirements	10
3.2.1	Sensor data collection (A)	10
3.2.2	Sensor data processing (B)	10
3.2.3	Reporting (C)	10
3.2.4	Data storage (D)	10
3.2.5	Estimating occupancy (E)	11
3.2.6	Data serving (F)	11
4	Design	12
4.1	High Level Summary	12
4.2	Components	12
4.2.1	Smart Sensor Boxes (A-C)	12
4.2.2	Centralised Storage (D)	14
4.2.3	HTTP Server (E-F)	16
5	Implementation	17
5.1	High Level Summary	17
5.2	Smart Sensor Boxes (A-C)	17
5.2.1	Sensor data collection	17
5.2.2	Sensor data processing	17
5.2.3	Sensor data reporting	18
5.3	Centralised Storage (D)	18
5.4	HTTP Server (E-F)	18
5.4.1	Estimating occupancy	19
5.4.2	Data serving	19
6	Evaluation	21
6.1	Sensor data collection	21
6.1.1	Test environment	21

6.1.2	Production environment	22
6.1.3	Non-functional Evaluation	23
6.2	Sesor data processing	23
6.2.1	Non-functional Evaluation	25
6.3	Reports	25
6.4	Centralised Storage	26
6.5	Estimating Occupancy	26
6.6	Data Serving	26
7	Conclusion	27
7.1	Summary	27
7.2	Future Work	27
7.3	Reflection	27
	Appendices	28

Chapter 1

Introduction

1.1 Context

The Internet of Things, or IoT is a term used to describe the expanding network of devices of many types, ranging from supercomputers to household appliances that contain embedded computer systems, making them capable of communication with other IoT devices. These devices aim to bring technology to every corner of our lives, bringing the idea of interconnected smart homes, workplaces, cities and more to reality.

Since the term was first used in 1999[2] the IoT has come a long way, with the number of internet connected devices rapidly increasing[1], and although nobody seems to agree on global market predictions, many predict massive growth[8][3][7].

One key area of study has been developing occupancy detection systems that would allow IoT devices to be more aware of the occupancy status of their surroundings. This is a crucial step to implementing smart homes, workplaces and campuses that are responsive and can meet their demands at any given time.

Advanced examples of this are already taking concrete form, with the University of Glasgow investing 1bn towards developing a smart campus[6] and the city of Barcelona making major overhauls to many public services in an attempt to bring the idea of a smart city to reality[9].

1.2 Motivations

The motivations for this project are wide, with smart sensors being applicable in a broad variety of ways. We will discuss the primary motivations for this project, which are understandably more relatable and accessible to myself, in addition to a number of secondary motivations that are considered due to close similarity.

1.2.1 Primary Motivation

To narrow our field of focus, we choose a main motivation to consider when producing the smart sensor boxes. We consider the following scenario: University administration decide they wish to upgrade buildings with smart sensor boxes with the overall goal being that students or other staff members can find occupancy information on certain areas remotely. For example library group study areas could be listed online with an indication of their current use, allowing students who plan on visiting to check availability beforehand. The University of Glasgow

library currently provides a similar feature, displaying computer use breakdown by floor, available from monitors within the library.

In addition the system should be built to be easily extendable, and capable of operating as a part of a larger IoT network. This is critical as there are many potential uses for such a system operating on a larger network of IoT devices. One example would be querying the system for occupancy information on a room and adjusting ventilation and heating based on the number of occupants, a effective system would bring potential for large savings in electricity and heating [4] while requiring very little human intervention or overhead costs.

1.2.2 Secondary Motivations

Although we focus on a university implementation of occupancy prediction, we can imagine a similar design could be incorporated into many different areas. Schools, workplaces and public buildings could all benefit from installing smart sensor functionality, allowing not only occupancy information to end users, but also providing more in-depth analysis of overall occupation to any interested administrators.

For example, smart public buildings could allow local councils to analyse the usage of various public services, allowing them to make more informed decisions on funding, maintenance, staffing and other key elements of operation.

To list a few more: Security systems, home-automation, public transport, gyms, dining locations and other areas of leisure. We can see that this problem is applicable to many areas of life and has the potential to shape cities and communities by providing accurate and reliable information that is used network-wide by many parties.

1.3 Goals

In short, the aim of this project is to design and construct low budget smart sensor boxes, capable of predicting human occupation in a room. To do this we use the popular brand of microprocessors Raspberry Pis in addition to a variety of sensors and other techniques.

In addition to the smart sensor boxes, we will design and implement a centralised hub capable of receiving, storing and serving information produced by a series of RPi boxes. This hub should also provide functionality for system-wide maintenance and configuration, to make setup and operation efficient and flexible.

This system should be built to be easily extendable, and capable of operating as part of the IoT network. This is critical as there are many potential uses for such a system operating on a larger network of IoT devices.

Chapter 2

Background Analysis

We start by researching and reviewing past studies conducted in similar fields, given that the subject matter is broad and widely applicable, we find many studies available. Many employ sensors and machine learning techniques to achieve meaningful results.

2.1 Human Detection

Sensors are a key element in almost every real-world computing application, especially those concerned with the detection of humans e.g automatic doors. Without them a computing device would have little to no information on its surroundings, and would be extremely ineffective for real-time interactive functionality. For this reason, sensors will be a primary component of the smart boxes.

Background research on previously conducted studies gives us a comprehensive list of applicable sensor, and their specific advantages and disadvantages.

2.1.1 Infrared Light

Infrared light is the electromagnetic radiation emitted by all objects with a temperature above 0K, giving the feeling of heat. The wavelength of IR light spans from 700 nm to 1 mm, existing just outside of the range of human vision[?]. PIR sensors feature a sensor face that can detect a change in the average amount of IR light hitting the surface, the sensor then outputs a signal indicating this change, with a low voltage indicating low change, and a higher voltage indicating a greater change[?].

They are most commonly used in PIR-based motion detectors such as automatic lighting systems. Generally the purpose of a PIR sensor is to detect the presence of humans at a binary level i.e true or false. However it has been shown that with suitable positioning and utilisation of machine learning techniques, a single PIR sensor can be used to predict the number of occupants in a room within a small range to fair accuracy[?].

There are many problems to overcome when using a PIR sensor for this purpose, most importantly, the positioning of the sensor; PIR sensors are prone to interference from occluding objects i.e blocking one another from the sensor. As a result of this, the most ideal position would be the ceiling of the room, however this may not prove as practical for other aspects such as maintenance[?].

Although it will be difficult to obtain a precise estimation by a PIR sensor alone, with the correct implementation, the sensor could potentially be used alongside other components to achieve a valid predictions.

2.1.2 Digital Images

The field of computer vision first emerged in the 1960s and has been developing steadily ever since with large amounts of research being done for a number of different problems. With endless applications from manufacturing to self-driving vehicles, its not surprising that there is a variety of methods available to us, specifically much research has been done on object detection and recognition.

Primitive solutions to object detection might include matching edges from input imagery to predefined templates or a taking and analysing a histogram of oriented gradients (HOG).

In recent years, more complex designs have been produced, often in combination with machine learning algorithms and a training dataset which has been shown to be effective in a variety of object detection problem areas.

Although more complex techniques are tempting, it is important to bear in mind the technical restraints put in place by using a microcontroller; computer vision techniques are notorious for being process intensive and many require the use of optimised graphical processors to achieve realistic processing times. Fortunately modern digital cameras are extremely advanced and cheap, high quality and compatible digital cameras are readily available, specifically the Raspberry Pi 3 Model B+ features a CSI camera port and a separate camera module. The camera module costs around 20, has an 8-megapixel sensor and is capable of recording in 1080p30 and 720p60[?] which is most definitely enough for our needs.

To conclude, it seems that computer vision could play a vital role in the final design and is definitely due consideration.

2.1.3 Audio

As humans, audio plays a crucial role in understanding our environment and communicating with each another, as a result we are able to detect human voices with almost absolute certainty. This close relationship suggests that it is a promising area of study in relation to human detection, i.e we use it successfully, therefore machine may be able to do the same.

We find that this is not quite as promising as previously thought, although a fair amount of research has been done on audio analysis by machine for high level tasks such as human detection, we have yet to see an implementation that can perform anywhere near the level of humans for such tasks. Most successful applications of audio analysis have been in speech recognition, for example speech to text functionality which has improved the accessibility of many computer systems and applications.

We find that there are a few key problems facing predicting human occupancy using audio, one of them being the fact that noise in the environment is constantly disturbing audio detections with noise, from a variety of sources. Much progress has been made over the past few decades concerning the analysis of a single human voice, however little has been made in developing solutions that would be applicable to mass human detection.

Although audio analysis clearly has the potential to be a major component in a smart sensor boxes, our findings suggest that the field needs a lot more development before audio analysis can give meaningful occupancy predictions.

2.1.4 Other

Luminosity

Luminosity sensors contain photodiodes which convert light to electric energy, allowing them to effectively measure the lighting level of there surrounding. The typical luminosity range of an office or classroom is between 0 and 500 lux, which is well within the standard sensor range of 0 to 40,000 lux . A commonplace example of an application of these sensors would be street lights which turn on and off. Not much information could be found on cases of luminosity sensors being used for human occupancy prediction, but we find that a cheap and easily available luminosity sensor would be able to provide an indication of occupancy, as a dark room would suggest that there are no occupants. We must also consider the possible interferences that may affect the luminosity of the room, or the reading produced by the sensor. Sunlight shining through a window onto the sensor would greatly increase its reading, and could give an inaccurate reading of the room lighting as a whole.

Temperature

Another sign of human occupancy could be room temperature, as occupants naturally heat the air through emission of body heat. However the standard precision of an appropriate temperature sensor is usually 0.5C, which it too low to be useful in determining exact temperature reading. The temperature change undergone when a group of people enter the room is unlikely be high enough to be properly detected by the sensor and be a guaranteed indication of human occupancy. This is especially true when the room temperature is subject to external factors, for example sun shining through the window will cause the room temperature to increase rapidly. We find that it would be too difficult to separate such an event from temperature rises that indicate human activity, so it follows that temperature sensors would be of little use for this problem.

CO2 and humidity

In an unregulated air system, human occupants would also give rise to CO2 and humidity levels through breathing, a higher number of occupants would see that the change in CO2 and humidity would be faster meaning the rate of change could give us a fairly accurate prediction. However for this very reason, building air must be properly regulated to ensure that it is safe for human occupancy, therefore any measurements we took would be expected to be within a fairly small range and would be highly influenced by external factors such as air conditioning, open windows and air flows.

2.2 Device Detection

Although the aim of this system is to detect human occupancy, everyday more and more people are using smartphones or similar electronic wireless devices for almost every aspect of life. It is estimated that 81% of adults in the UK own a smartphone (as of 2016)[5], and with most of these people using their devices everywhere they go, it is reasonable to say that the number of wireless devices in a specific area would be very highly correlated with the number of human occupants. We must also consider the possibility of rooms containing wireless devices that are not relative to the human occupants, for example lab computers may contain Bluetooth functionality that would distort the human to device ratio, however typical workplace machines operate using wired communication, as it is more practical for a stationary machine.

Over the years, smartphone development has brought about faster and more robust wireless protocols and methods. Since many of these protocols are the same ones used by microcontrollers, it is reasonable to think we may also be able to get an estimation of wireless devices nearby using similar techniques to human detection.

Although this system would greatly benefit from being able to detect all devices in the nearby area, obviously this is not technically feasible with any device, since many users would choose to not broadcast their devices and many protocols are not built with device broadcasting in mind, for example communicating on 3G will not let all other nearby 3G devices know the location such as bluetooth communication might.

The only feasible device detection method we can see is using Bluetooth and Bluetooth Low Energy device scanning, which will receive alerts from nearby devices that have broadcasting enabled. Although not every bluetooth device will broadcast its location in response to a device scan, the ratio of devices that respond should be relatively constant for similar audiences.

2.3 Existing products

During our background analysis we find that similar systems have been implemented before for real-world production use.

"Meshium", a product of the IoT provider Libelium, attempts to detect WiFi and Bluetooth communications in the nearby area to obtain a number of devices. It is estimated to detect up to 95% of nearby wireless devices operating on these protocols[?].

Another IoT solutions provider "Retail Sensing" have developed a CCTV based people counting system which utilises computer vision techniques with CCTV camera recordings from cities, shopping centres etc. to predict human occupancy by area to up to 98%[?].

Unfortunately since these systems are a product costing money, the intellectual property is protected and technical specifications are vague and mostly refer to "our processing algorithms". However we can at least see that these products are feasible and obtain a general overview of how they work which will be a useful reference during the design stage.

Chapter 3

Requirements

3.1 Functional Requirements

We begin by defining 6 functional requirements (A-F) for the final system and their associated sub-requirements. The order of requirements should give an indication of the data flow, starting at the sensors and ending with the end user.

3.1.1 Sensor data collection (A)

The smart sensor boxes should be capable of collecting data from connected sensors for a specified period of time. The connected sensors we should support are:

1. Bluetooth device scanning.
2. Bluetooth Low Energy device scanning.
3. Camera (still images)

3.1.2 Sensor data processing (B)

In addition to data collection, the smart sensor boxes should be capable of performing some form of post-processing on the sensor data. For each implemented sensor, we choose an appropriate processing technique that can be applied:

1. Bluetooth and Bluetooth LE devices.
 - (a) Log any bluetooth responses received to file.
 - (b) Anonymise device addresses for ethical reasons (only if required).
 - (c) Collate results from both sensors.
 - (d) Remove duplicate results where possible.
 - (e) Retain results for a given period after detection.
 - (f) Attempt to smooth out connection strength values for each device.

2. Images

- (a) Log any images taken to file.
- (b) Apply an object recognition algorithm to the still image.
- (c) Log any object recognition output images to file.
- (d) Parse output for number of humans found.

3.1.3 Reporting (C)

The smart sensor boxes should be able to:

- 1. Smooth the output of camera post-processing results to obtain an more reliable metric for a given time.
- 2. Collect all devices from device post-processing that are still valid (within decay time range).
- 3. Compile a report from: the smoothed people number, the number of devices, the current time, any report meta-data required.
- 4. Send this report securely to an external component for further use.
- 5. Log any reports that are given to file.

3.1.4 Data storage (D)

The system must be capable of storing data about:

- 1. Buildings, floors, rooms.
 - (a) Name.
 - (b) Description.
- 2. Smart boxes.
 - (a) Name.
 - (b) Description.
 - (c) Location.
 - (d) Authentication data.
- 3. Sensor output.
 - (a) Time.
 - (b) Sensor reading.
- 4. Smart box reports.
 - (a) Time.
 - (b) Sensor reading.
- 5. Estimations.
 - (a) Time.

- (b) Location.
 - (c) Estimate.
- 6. Manual readings
 - (a) Time frame.
 - (b) Location.
 - (c) Reading.
- 7. Admin users.
 - (a) Username.
 - (b) Hashed password (it is unsecure to store plain passwords for many reasons).

3.1.5 Estimating occupancy (E)

The system should be able to use smart box reports, to give an estimation of occupancy for each location with reports.

1. Generate some form of regression model for each sensor type.
2. Train each model using a combination of smart box reports values and true readings for the associated room and time range.
3. We should be able to continue to train the models on and off, so they can continually improve throughout service.
4. Use new reports to generate occupancy estimates using an average of all recent report data for each smart sensor box.
5. The system should be capable of making estimated with missing report values to ensure that not all sensors are essential for every smart box node.
6. The system should also be able to use reports from multiple smart sensor boxes in the same area, and take an average of the seperate estimations.

3.1.6 Data serving (F)

1. We should provide public available REST API endpoints that allow access to data on:
 - (a) Buildings.
 - (b) Floors.
 - (c) Rooms.
 - (d) Room estimates.
2. We should provide private (user only) available REST API endpoints that allow addition, removal and modification of data on:
 - (a) Buildings.
 - (b) Floors.
 - (c) Rooms.

- (d) Smart boxes.
 - (e) Users.
3. We should provide a front-end display that provides an interface for the public api to end users.
 4. We should also provide a front-end display that provides an interface for the private api.

3.2 Non-functional Requirements

For each functional requirement we also state several non-functional requirements that are still essential to an effective solution.

3.2.1 Sensor data collection (A)

- Ideally each smart sensor box should be as cheap as possible while providing adequate data, in order for the system to be scaled up without large overhead costs.
- If the data cannot be collected continuously, meaning they require some amount of time to compute results before obtaining a new reading, then we should perform as many iterations as possible to ensure that results are up to date.
- The data collection should be able to run for as long as possible without human intervention to reduce the amount of maintenance required.

3.2.2 Sensor data processing (B)

- The post processing should be as quick as possible to ensure that final reports are not heavily delayed.
- The processing required should be as light as possible, as the Raspberry Pi does not have a large processing capacity. We should also seek to use optimised processing implementations where possible.

3.2.3 Reporting (C)

- To ensure that occupancy reports and predictions are not outdated, RPIs should be able to send reports at a rate of at least one report per minute. The reports should also stick to the specified rate with fair accuracy and not deviate too far from a steady report rate.
- Extensibility - the report formatting should take a standard form such as JSON and be easily extendible to allow the use of more sensors.

3.2.4 Data storage (D)

Any data stored should be:

- Implemented in an efficient and correct manner, without unnecessary dependencies between data.
- Scalable without any major modifications.

- Accessible as quickly as possible and with no down-time.
- Secure from malicious attack or unauthorised access.

3.2.5 Estimating occupancy (E)

- The estimates should be as accurate as possible.
- Estimates should be produced at a reasonable enough rate to ensure that they are not out of date.

3.2.6 Data serving (F)

Any data stored should be:

- User friendly and easy to use.
- As up to date as possible.

Chapter 4

Design

4.1 High Level Summary

We start by splitting the requirements to three separate components with individual functions.

1. Smart Sensor Boxes: Collects data from the various sensors, also generate reports detailing RPi identification, along with smoothed sensor data. Reports are then sent to the HTTP server for analysis.
2. Centralised Storage: A database that allows for the storing and querying of data must contain information on building, RPis, estimates, reports, users etc.
3. HTTP Server: A simple HTTP application that can receive sensor report data from the smart sensor boxes. The sever should also be able to perform estimations for locations based on reports and also feature a web application that serves this data to the users.

Figure 4.1 Shows a simple high level diagram of the system design.

4.2 Components

4.2.1 Smart Sensor Boxes (A-C)

The smart sensor box logic consist of a main reporter algorithm which initialises scanning for each detection method and monitors the output. Every time a report is required, the algorithm takes the outputs from each output queue, and collates them to a report which is sent to the HTTP server for furthur use. Algorithm 3 shows the pseudocode for this algorithm.

The detection methods supported are:

1. Bluetooth / BTLE device scanning - Continuously scans for devices and adds any discoveries to a output map, scans are done for a specified time before restarting to ensure that devices do not stop responding the scan signal. Algorithm 1 details the algorithms used.

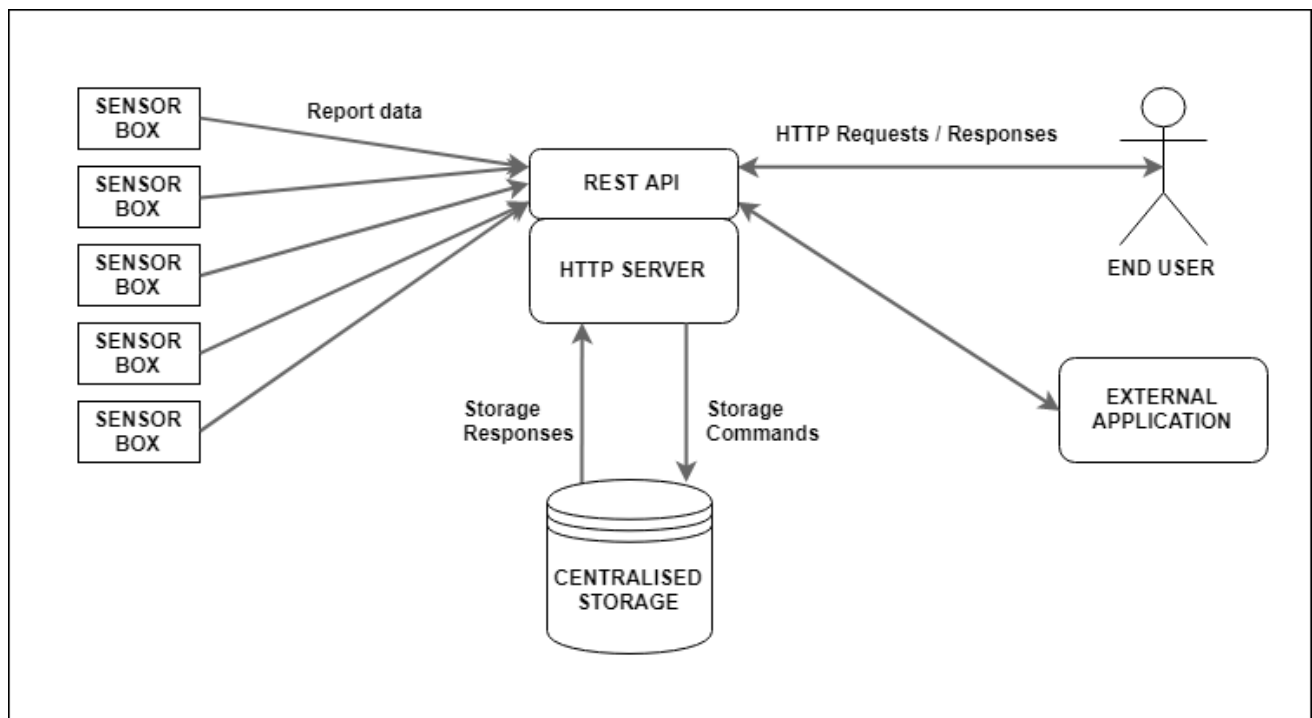


Figure 4.1: High level design of system.




Figure 4.2: High level diagram of the smart sensor box implementation.

2. Camera and object detection - Runs on a loop taking images and feeding them into an object detection algorithm which predicts objects and their certainty. The number of people is extracted from the detection output and smoothed with previous results using an average of the previous value and the new prediction. Algorithm 2 details the algorithms used.

```

1 void scanBluetooth(Int  $T_{cycle}$ , Int  $T_{timeout}$ , Int  $T_{decay}$ , Map output)
2 begin
3   while  $T_{timeout} > T_{now}$  do
4     scan( $T_{cycle}$ ,  $T_{decay}$ )
5     while scanning do
6       discovery  $\leftarrow$  new_discovery
7       handleDiscovery( $T_{decay}$ , discovery, output)
8     end
9   end
10 end

8 void handleDiscovery(Int  $T_{decay}$ , Map discovery, Map output)
9 begin
10  addr  $\leftarrow$  discovery['address']
11  rssi  $\leftarrow$  discovery['rssi']
12  time  $\leftarrow$  discovery['time']
13  if address in output then
14    old  $\leftarrow$  output[addr]
15    if old['time'] +  $T_{decay} < T_{now}$  then
16      new  $\leftarrow$  {'rssi' : (rssi + old['rssi'])/2, 'time' : time}
17      output[addr]  $\leftarrow$  new
18    else
19      new  $\leftarrow$  {'rssi' : rssi, 'time' : time}
20      output[addr]  $\leftarrow$  new
21    end
22  else
23    new  $\leftarrow$  {'rssi' : rssi, 'time' : time}
24    output[addr]  $\leftarrow$  new
25  end
26 end

```

Algorithm 1: Pseudocode for Bluetooth and Bluetooth Low Energy device scanners.

4.2.2 Centralised Storage (D)

The centralised storage component should take the form of a database. We start by defining the entities as:

- Building
- Floor
- Room
- Rpi

```

1 void scanCamera(Int  $T_{cycle}$ , Int  $T_{timeout}$ , Int output)
2 begin
3   while  $T_{timeout} > T_{now}$  do
4      $img \leftarrow takeImage()$ 
5      $objs \leftarrow objectDetection(img)$ 
6      $hum \leftarrow numHuman(objs)$ 
7      $output \leftarrow (output + hum)/2$ 
8      $wait(T_{cycle})$ 
   end
end

```

Algorithm 2: Pseudocode for the camera detection algorithm.

```

1 void reporter(Int  $T_{cycle}$ , Int  $T_{timeout}$ )
2 begin
3    $T_{decay}, T_{ccycle}, T_{bcycle} \leftarrow getConfig()$ 
4    $auth \leftarrow getAuthData()$ 
5    $bt\_output \leftarrow new \textbf{Map}$ 
6    $cm\_output \leftarrow 0$ 
7    $scanBluetooth(T_{bcycle}, T_{timeout}, T_{decay}, bt\_output)$ 
8    $scanCamera(T_{ccycle}, T_{timeout}, cm\_output)$ 
9   while  $T_{timeout} > T_{now}$  do
10     $devices \leftarrow bt\_output$ 
11     $people \leftarrow cm\_output$ 
12     $report \leftarrow \{'devices' : devices, 'people' : people, 'time' : T_{now}, 'auth' : auth\}$ 
13     $sendReport(report)$ 
14     $wait(T_{cycle})$ 
  end
end

```

Algorithm 3: Pseudocode for the reporter algorithm.

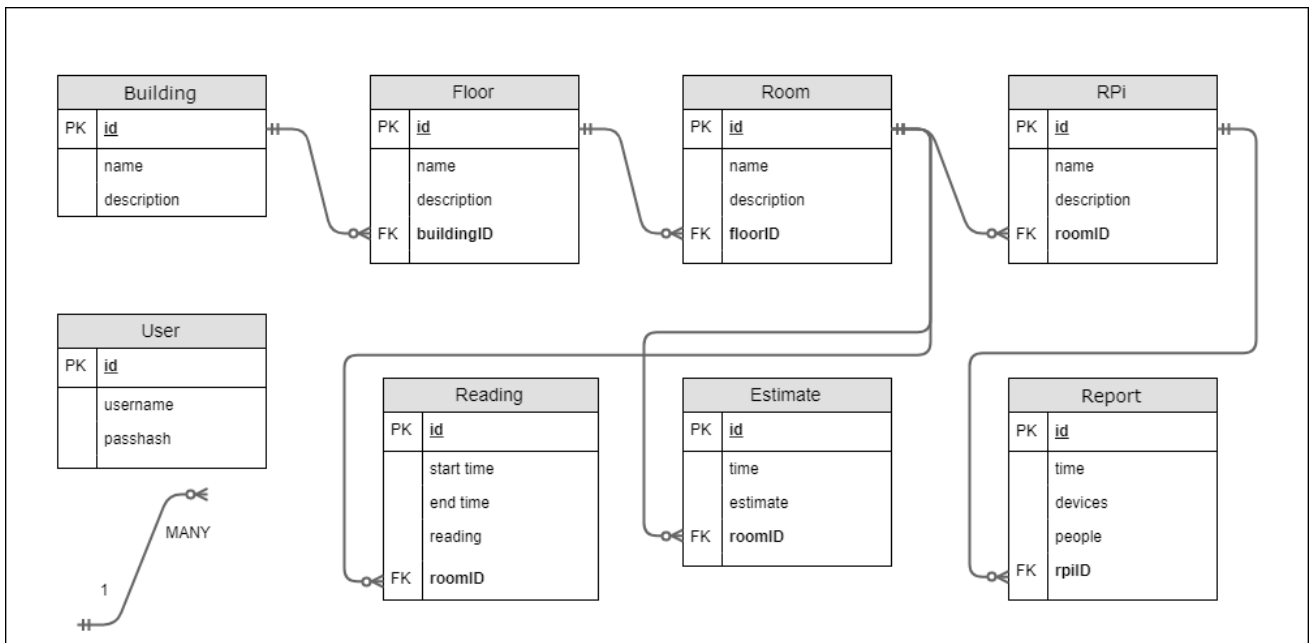


Figure 4.3: ER diagram of database design.

- Report
- Estimate
- Reading
- User

In addition to the data fields set out in the requirements stage, we must use foreign keys to map out the relationships between these entities:

- Buildings may contain one or more floors.
- Floors may contain one or more rooms.
- Rooms may contain one or more RPIs.
- RPIs may produce one or more reports.
- An estimate has one associated room, and room may have one or more estimates.
- An reading has one associated room, and room may have one or more reading.

To ensure that each entity instance can be uniquely identified, the database should store an associated integer id for each table row. We must also store an authorisation key for each smart sensor box, that will allow the server to verify the validity of any incoming reports. It is also important to note that when storing passwords we should hash them beforehand, this ensures that passwords are not compromised, while still allowing us to verify user logins.

Figure 4.3 Shows an ER diagram of database design.

4.2.3 HTTP Server (E-F)

Chapter 5

Implementation

5.1 High Level Summary

5.2 Smart Sensor Boxes (A-C)

The smart sensor boxes were implemented using the popular brand of microprocessors Raspberry Pi, the RPi were a great choice for a smart sensor box such as this due to the fact that it is compatible with a wide variety of sensors and have relatively powerful computation ability. Specifically we use the most recent version, Raspberry Pi 3 Model B which features inbuilt WiFi and Bluetooth adapters, as well as a dedicated RPi camera port.

5.2.1 Sensor data collection

Since the RPi 3B has an inbuilt Bluetooth adapter, we can use the Python module `bluez` to constantly perform a scan for nearby Bluetooth devices that are broadcasting nearby. We find an example Python script with similar desired functionality, we modified this file by adding a 'start' function that other Python modules can call with an output queue. Similarly, we use the module `bluepy` to scan for Bluetooth Low Energy devices, implementation is relatively simple, we perform a standard BTLE scan with a custom class to handle devices, which are sent to the BTLE output queue.

We collect still images from the camera module using the Python module "picamera" which provides a Python interface for the Raspberry Pi camera module. Obtaining an image is as simple as a function call, with a filename as argument.

5.2.2 Sensor data processing

To process all Bluetooth and BTLE devices, we create a thread-safe dictionary object that holds device information as value, with device address as the key. When we perform device scanning, we pass a separate output queue to each device scanner and periodically check for new additions, these additions are taken from the queue and inserted to the dictionary using a custom 'add' method.

Captured image files are then immediately passed by argument through a terminal command to an implementation of the object detection algorithm You Only Look Once also known as YOLO, once the detection algorithm

has completed, the standard output can be parsed for detected objects and their associated probability. We then count the occurrences of "person" to get a integer reading, which is added to the output queue.

An original implementation darknet was chosen, however the RPi inbuilt GPU is not very powerful and was taking around 180 seconds per still image. To mediate this we switched to a darknet version that was optimised for ARM processors, which brought the total time taken down to 35 seconds. This would have been sufficient, however it was found that by using a less intensive configuration files package, tiny-YOLO, we could improve the time further to 10 seconds, with relatively little difference in accuracy.

5.2.3 Sensor data reporting

Reporting the processed sensor data is relatively simple, we built a Python script that initialises all three scanning methods with seperate output queues

This functionality is implemented as 3 separate python scripts that can be called as subprocess from the reporter process. We pass in an output queue to each script on execution to allow the main process to read any devices and YOLO predictions. 4.2

5.3 Centralised Storage (D)

We implement our centralised storage using a Google Cloud SQL instance which hosts a MySQL 5.7 server containing our databases. The users table is implemented as one database named "users", all the other tables are contained within another database named "reports". To access the SQL instance, we add the IP address of the connecting device to the authorised networks list. To ensure that all data transmitted to and from the SQL instance is encrypted end to end, we can enforce SSL connections and generate client certificates that can be used by client devices (servers, development environments etc.).

For added protection against database corruption or loss of data, we can enable automatic backups through the Google Cloud console which would allow us to restore lost data from previous copies of the database.

5.4 HTTP Server (E-F)

In order to quickly develop and deploy a web server, we use the Python microframework Flask which requires very little setup or configuration.

To access the central storage instance, we create a user for the server and add the IP to authorised network list. Then we can use the Python module "MySQLdb" which provides an interface for Python to use MySQL. In order to keep the communication secure we must also store the client SSL certificates with the application so that they are accessible to MySQLdb.

Although suitable for development of web applications, Flask alone is not sufficient for deploying applications to a production environment due to the fact that it scales poorly. To use our Flask application in a production environment, we need to use "mod_wsgi" package which will host our application, making it suitable for realistic usage.

The production environment is installed on a Google Cloud compute engine which serves the webpage, to serve this to the public we obtain the domain "peoplecount.cf" and route DNS for that domain to the static IP

address of the instance.

Since the data our application will be sending and receiving must be secure, we must ensure that our application serves and receives data through HTTP over TLS (HTTPS), which will ensure that the data is end-to-end encrypted and safe from interception. To make sure that HTTPS is enforced, we set up a redirect from any HTTP requests to the HTTPS equivalent.

5.4.1 Estimating occupancy

We choose to use the Python machine learning module "sklearn", since it has a wide range of machine learning techniques to choose from and is very easy to use. Specifically we choose to use the Stochastic Gradient Descent Regressor model (SGDRegressor), which trains one sample at a time, updating each time at a decreasing learning rate. This allows us to continually train the model with new training samples, while still making estimates on real data. This type of regression is best suited for large amount of training data, which is good as it allows us to keep training, if the model is not accurate enough.

We implement a separate regression model for each type of report value, devices and people seen. This means that the system can still make estimates for rooms that are not fully equipped or are faulty. Reports are split into their separate values, with regression being performed on each non-null value, a final estimate is taken by averaging the separate results. In addition it makes the regression process a little more extensible, new sensor types could be added without updating existing smart sensor hardware or software.

We schedule the whole estimation process to happen every 60 seconds, this functionality is implemented using 'flask_apscheduler' which provides an advanced process scheduler. We simply set a job configuration with a 60 second timer, and start the scheduler at runtime.

Since we need our regression models to persist, even if the application terminates, we can use the Python 'pickle' module, which can dump the regression instances to a pickle file. When the estimates are being performed, the models can be loaded from file and used before discarding the new model instances. Similarly we can load a model, continue training with new training data, and then overwrite the pickle file with the updated model.

5.4.2 Data serving

Web pages are served through the url scheme "peoplecount.cf/webapp/...", when a web page is requested, the application will send the associated HTML templates, however the data is populated by server side JavaScript using asynchronous requests to the public and private API endpoints. This allows us to simplify the web page response code while maintaining functionality through the existing API. In addition, it allows us to set repeated timers on these requests, so the page is constantly being updated with the most recent data.

The server must also allow admin functionality, this requires us to use some method of authentication for any admin users. We use the Flask-login builtin functionality to implement a simple username/password system with new user registration restricted to admin users. Endpoints that require authentication can now simply be wrapped with "@login_required" which will ensure that any responses not authenticated will be required to provide login details. After logging in, authentication data is stored in session cookies that are safe from tamper.

In order to ensure the GUI is intuitive, we use familiar icons in place of text / buttons to minimise the learning curve. For example a '+' icon is immediately associated with addition, similarly a waste bin icon will be associated with removal or deletion. Fortunately, we can use 'Font Awesome Web Application Icons' which feature a wide range of intuitive icons for free commercial use. Furthermore, we use colour to give a general

indication of occupancy when viewing a building or floor page. For example a busy room would have a red background, compared to a mostly empty room which would have a green one.

All API endpoints follow the scheme "peoplecount.cf/api/...". An example of both of these would be "peoplecount.cf/webapp/home" and "peoplecount.cf/api/buildings/get-all".

Chapter 6

Evaluation

We perform some simple test on the Python modules produced to ensure that they work as planned, however the codebase could have benifited greatly from using automated testing. Unfortunately this was not considered until a later date and as a result the code is difficult to effectively test without either major refactoring or extensive use of mock components. In terms of the smart sensor boxes, the algorithms for data post-processing and reporting could benifit from being rewritten to be highly modular, which would allow for easy unit testing. In addition, the public and private REST API could be more thoroughly evaluated if integration test were used with a mock database which allowed the server to test all functionality without damaging any of the production data. In retrospect, this is a key area of improvement.

6.1 Sensor data collection

We evaluate the sensor data collection methods in two environments, a test environment (my home) and a production scenario where smart sensor boxes are placed in the University of Glasgow Boyd Orr building.

6.1.1 Test environment

We attempt to evaluate the types of devices found for both device scanning methods, however only a small number of devices were personally available. Tests were done by performing a background scan to collect the idle state of the test environment, then tests were conducted one by one while monitoring scanner output. We find that both scans have their seperate strengths and weaknesses:

We found that it was possible for a standard bluetooth device to be detected using a Bluetooth scan upto a distance of 10-13m, even if the devices were in another room. However we should note that walls in university buildings are generally a lot thicker and are likely to be composed of materials with a higher blocking power than residential buildings. We also found that some mobile devices sometimes had to be unlocked and on a settings page before they would broadcast presence. In addition we were unable to detect smaller devices such as a wireless mouse. For Bluetooth devices., we found that MAC addresses appeared to be genuine i.e devices did not broadcast using a anonymous address.

In contrast, the Bluetooth Low Energy scans did not discover mobile phones, but proved more sucessful at detecting smaller Bluetooth devices. We also found a large number of unidentified devices, since all BTLE devices inside my flat were accounted for, it is reasonable to assume that Bluetooth Low Energy signals are able to transmit more effectively through walls and obstructions. We also found that laptops were also discoverable,

Device	Bluetooth	Bluetooth LE
Smart phone (1)	Found with static address when on Bluetooth settings screen.	Not found.
Smart phone (2)	Found with static address..	Not found.
Laptop (1)	Found with static address.	Found with changing address.
Laptop (2)	Found with static address.	Found with changing address.
Tablet	Found with static address when on Bluetooth settings screen.	Not found
Wireless mouse (1)	Not found	Found with static address.
Wireless mouse (2)	Not found	Found with static address.

Table 6.1: BT and BTLE scan home test findings.

however restarting the Bluetooth service would change the address broadcasted. Table 6.1 gives a breakdown of devices used and the findings.

In addition we can test the image detection, this is relatively simple as we can just take some images, and access them through VNC Viewer connection to confirm they are correct. Appendix ?? contains some example images that have been captured

6.1.2 Production environment

As previously stated, we also deploy two Raspberry Pi smart sensor boxes to the UoG Boyd Orr building. We choose rooms BO720 and BO715 for the following reasons:

- Both rooms are equipped with suitable infrastructure (power and ethernet).
- Both rooms are accessible to myself.
- Rooms are of similar size and are therefore more easily comparable.

Due to ethical considerations, it would not have been reasonable to utilise the camera sensor within the lab, as occupants would need to give prior consent which is highly unfeasible to pursue. However we were able to perform a production evaluation for the Bluetooth and Bluetooth Low Energy scanners, given that we address the ethical concerns:

- Device address collection - Although the smart sensor boxes briefly store device MAC addresses in memory, they are immediately hashed to assure anonymity.
- Transparency - Smart sensor boxes were situated in plain sight with clear explanation and contact information. Figure 6.1 shows the setup in BO720.

Our findings are unexpected, however they provide insight into the issues associated with a use in a production environment that were not uncovered during the first stage of evaluation. Most notably we found that Bluetooth Low Energy devices detection made up around 90-100% of the total devices, with standard Bluetooth barely ever being detected.

** DEVICES / OCCUPANTS GRAPH HERE




Figure 6.1: The production evaluation setup in BO720.

6.1.3 Non-functional Evaluation

To fully evaluate the smart sensor boxes, we also evaluate the status of each non-functional requirement set out in the requirements chapter.

- Price - We calculate the standard price of a single smart sensor box to be around 52, however purchasing in bulk would reduce this cost considerably, appendix ?? shows the calculations made.
- Data collection rate - Both BT and BTLE scans are continuous which is ideal, the image capturing however, requires some processing time. We take the average time required to take an image using 5 separate captures, we find that the time varies very little and takes an average of 4 seconds per image, which is definitely fast enough for our needs.
- Maintenance required - During the production environment tests, we found that errors in the scanning processes would crash that scan, and the reporter must be manually restarted in order to restart the scanning. However the code could be modified slightly to log any error to file and attempt an automatic restart.

6.2 Sensor data processing

We perform a set of manual tests to gain an evaluation of the post processing algorithms. Tables 6.2 and 6.3 detail the manual tests performed, we find that the algorithms are mostly sufficient, however there are several opportunities for further improvements. Additionally we encountered some limitations on the system that were unavoidable as far as my knowledge.

Sub-requirement	Test	Result	Status
Logging	Manually compare standard output with log files.	We find that devices are being correctly logged.	Adequate.
Anonymisation of addresses	Monitor output of scan with anonymisation set as true.	Device addresses are successfully one-way hashed, hashing the same address produces the same result.	Adequate.
Collation of results	Monitor output of both threads, and of parent thread.	Devices are merged and parent output contains all devices found.	Adequate.
Duplicate removal	Monitor output of both threads, and of parent thread.	We find that when devices broadcast an anonymous address or respond to both Bluetooth and BTLE scans with different addresses, there is no way to merge the discoveries as they appear to be separate devices.	Limitations.
Keeping results	Output is monitored with controlled devices.	Devices data persists for the required amount of time and disappears afterwards.	Adequate.
Smoothing results	Manual calculations compared to post-processing standard output.	Device RSSI smoothes according to the weighted average definition. However the weightings could potentially be optimised for more accurate readings.	Possible improvements.

Table 6.2: BT and BTLE device processing test findings.

Sub-requirement	Test	Result	Status
Logging	Verify that any captured or processed images are saved to file.	We find that all images are being correctly logged.	Adequate.
Object recognition	Basic test with test image set and home-use, manually comparing bounding boxes with actual object in the image.	We find that there are many situations where people are not recognised e.g occlusion from other objects or bad angles. We could continue to train the object recognition model to obtain as accurate results as possible.	Possible improvements.
Parsing	Monitor output of object detection , and of parsing thread.	As expected, this is a simple parsing operation and has no issues.	Adequate.
Smoothing results	Manual calculations compared to post-processing standard output.	The number of people smooths accordingly, however the weightings could potentially be optimised for more accurate readings.	Possible improvements.

Table 6.3: Image processing test findings.

6.2.1 Non-functional Evaluation

As before, evaluate the status of each non-functional requirement.

- Data processing speed - Since devices are loaded into the output object in virtually real-time, we must only consider the image post processing. We found that over 5 cases, the average time taken to post-process the image was 10 seconds.
- Data processing intensity - We find the same is true with process intensity, the device processing is negligible, compared to the image processing which is very processor intensive, however I believe we have one of the best suited YOLO implementations and the nature of image processing tasks is unavoidable.

6.3 Reports

Additionally we perform a set of manual tests on the reporter script:

Sub-requirement	Test	Result	Status
Collection of post-process state	Verify that any captured or processed images are saved to file.	We find that all images are being correctly logged.	Adequate.
Reporting	Verify on sever side that reports are being recieved and have correct structure.	We find that reports that have no associated camera, give people = 0 instead of people = null. Additionally server side errors i.e 500 Server Error, will crash the reporter script upon response.	Minor issues.
Logging	Monitor output of reporter script, and compare to log file.	We find that log files match correctly.	Adequate.

6.4 Centralised Storage

6.5 Estimating Occupancy

6.6 Data Serving

Chapter 7

Conclusion

7.1 Summary

7.2 Future Work

7.3 Reflection

Appendices

Bibliography

- [1]
- [2] Kevin Ashton. That 'internet of things' thing. *RFID Journal*, 2009.
- [3] Peter Bowen, Asit Goel, Michael Schallehn, and Michael Schertler. Bain insights. *Bain Insights*, 2017.
- [4] Michael J. Brandemuehl and James E. Braun. The impact of demand-controlled and economizer ventilation strategies on energy use in buildings. *ASHRAE Transactions*, 1999.
- [5] Deloitte. Mobile consumer 2016 report, 2016.
- [6] Kristian Hentschel, Dejice Jacob, Jeremy Singer, and Matthew Chalmers. Supersensors: Raspberry pi devices for smart campus infrastructure (short paper). *2016 IEEE 4th International Conference on Future Internet of Things and Cloud (FiCloud)*, 2016.
- [7] IDC. Worldwide internet of things forecast 20152020. 2015.
- [8] GrowthEnabler & MarketsandMarkets. Market pulse report. Technical report, 2016.
- [9] Ross Tieman. Barcelona: smart city revolution in progress. *Financial Times*, 2017.